MIDTERM

INFO 7390: Advances in Data Sciences/Architecture *Team 9*

Bhavesh Patel Ila Nigam Ayushi Srivastava

Table of Contents

1. Problem1

- Goal
- Power BI and Inferences
- Logistic Regression
- Neural Network
- Classification Tree
- Discussion

2. Problem2

- Goal
- Data Cleaning
- Power BI and Inferences
- Logistic Regression
- Neural Network
- Classification Tree
- Discussion

3. Problem3

- Goal
- Power BI and Inferences
- Linear Regression
- Neural Network
- Regression Tree
- Discussion

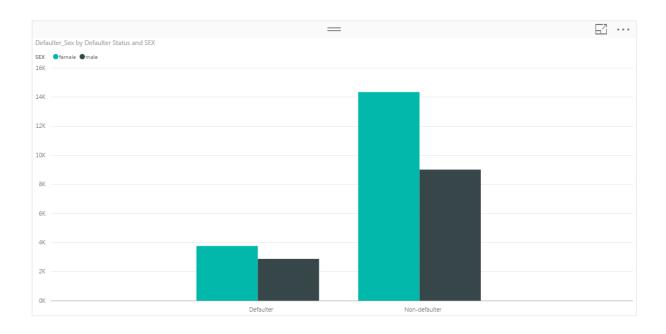
Problem1- Credit Card Clients

♣ GOAL

- 1. The data set contains instances and 23 attributes consisting of gender, education profile, marital status, age, history of statement balance, payment status and binary status of default (1 or 0).
- 2. The goal of the problem is to clean the data and infer various relations between the predictor variables. Using logistic regression, neural network and classification tree we have to classify the data and predict the values.
- 3. We have to build different models to evaluate different trends and compare the sensitivity, Specification and plot confusion matrix, ROC Curve and Lift Charts for all the classification method.

♣Power PI and Inferences

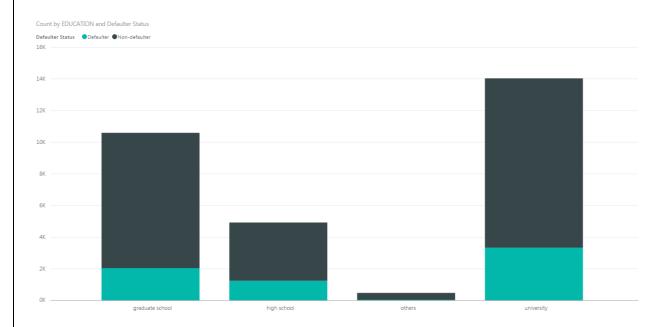
Gender profile of defaults payment vs Non-Default



Gender profile of defaults payment vs Non-Default

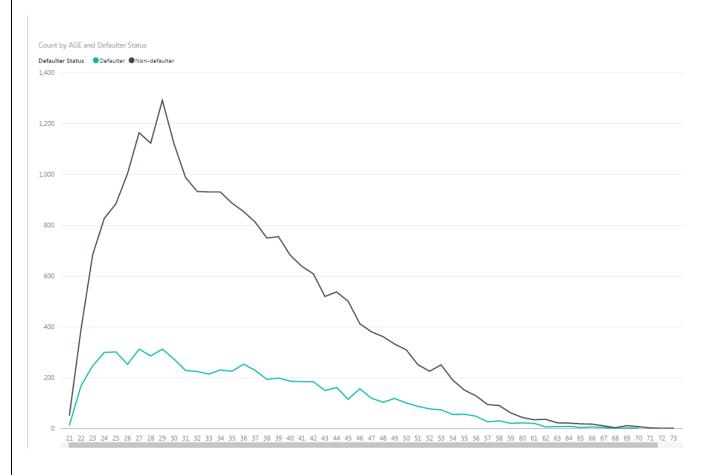
✓ We can see from the above bar chart, number of females is higher than males for both default and no default payment categories, but we cannot see their respective

Education profile of defaults payment vs Non-Default



✓ We can see from above bar graph, ratio of customers having graduate school
is higher in default payment category as compared to no-default and ratio of
customers having university education is higher in No Default Payment than
Default category

Density vs. Age Profile



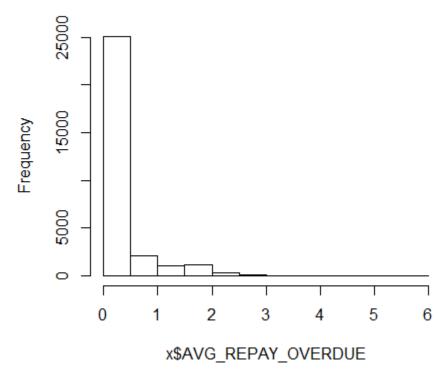
✓ We can see clearly see that age peaks around 28-29 years in default category and it has lower peak around 27-28 years in no default category.

Logistic Regression

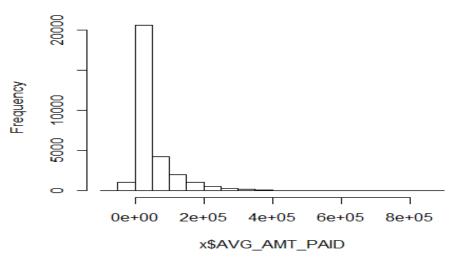
- 1. To go ahead with logistic regression model we first cleaned the data. Removed the outliers. Plotted various graphs such as Histograms, Box plot to maintain the integrity of data.
- 2. Next we sampled the data into Test and Train to model the train data and then evaluate our model on based of Test data set.
- 3. In Logistic regression we did GLM modeling with variables LIMIT_BAL+SEX+MARRIAGE+AVG_REPAY_OVERDUE
- 4. Threshold of 0.27 was set to classify data into 0 or 1 and predictions were made as per the analysis.
- 5. Graphs were plotted to evaluate the performance of the models

Predictor analysis

Histogram of x\$AVG_REPAY_OVERDUE



Histogram of x\$AVG_AMT_PAID



```
#Sampling the data
smp_size <- floor(0.75 * nrow(x))
set.seed(123)
train_ind <- sample(seq_len(nrow(x)), size = smp_size)

train <- x[train_ind, ]
test <- x[-train_ind, ]
names(test)
#Build Logistic Regression
glm.fit <- glm(Y~LIMIT_BAL+SEX+MARRIAGE+AVG_REPAY_OVERDUE,data=train, family=binomial(link="logit" summary(glm.fit)

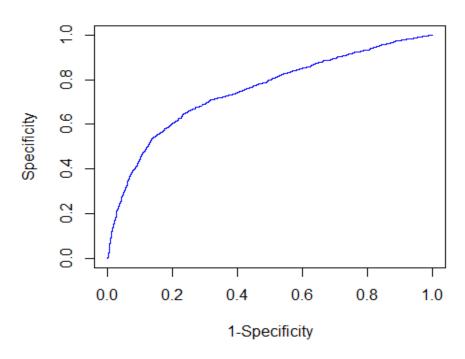
prob <- predict(glm.fit, newdata=test, type="response")
pred <- rep(1,length(prob))
pred[prob<=0.27] <- 0</pre>
```

```
glm(formula = Y \sim LIMIT_BAL + SEX + MARRIAGE + AVG_REPAY_OVERDUE +
    AVG_BILL_REPAY + AVG_AMT_PAID, family = binomial(link = "logit"),
    data = train)
Deviance Residuals:
                   Median
   Min
              10
                                          Max
-3.3625 -0.6358 -0.5635 -0.3821
                                       3.3093
Coefficients:
                   Estimate Std. Error z value Pr(>|z|) -9.269e-01 8.052e-02 -11.511 < 2e-16 ***
(Intercept)
                                           -8.118 4.75e-16 ***
-4.537 5.71e-06 ***
                               1.564e-07
                   -1.270e-06
LIMIT_BAL
SEX
                   -1.511e-01
                               3.330e-02
                                           -5.238 1.63e-07 ***
MARRIAGE
                   -1.653e-01
                               3.156e-02
AVG_REPAY_OVERDUE 1.278e+00
                               2.910e-02
                                           43.941 < 2e-16 ***
                                                   < 2e-16 ***
                   -3.215e-05
                               3.552e-06
                                           -9.051
AVG BILL REPAY
                                            3.826 0.00013 ***
                   1.201e-06 3.138e-07
AVG_AMT_PAID
Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' ' 1
(Dispersion parameter for binomial family taken to be 1)
    Null deviance: 27008 on 25499 degrees of freedom
Residual deviance: 23551 on 25493 degrees of freedom
AIC: 23565
Number of Fisher Scoring iterations: 5
```

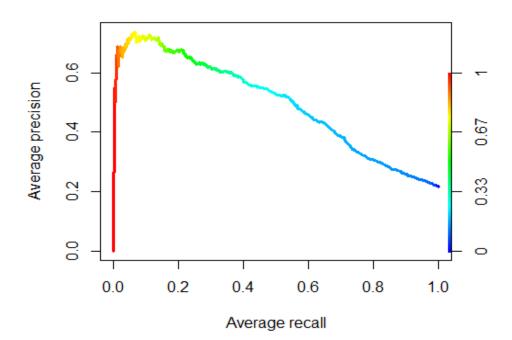
```
> #Plotting ROC Curve
  > #install.packages("ROCR")
  > library(ROCR)
  > prediction <- prediction(prob,test$Y)</pre>
  > performance <- performance(prediction, measu
 > plot(performance, main="ROC curve",xlab="1-5
> # ROC area under the curve
  > prediction <- prediction(prob,test$Y)</pre>
  > auc.tmp <- performance(prediction, "auc")</pre>
  > auc <- as.numeric(auc.tmp@y.values)</pre>
 > print(auc)
  [1] 0.7521183
confusionMatrix(pred, test$Y)
Confusion Matrix and Statistics
          Reference
rediction
            0 1
         0 3177 541
         1 351 431
               Accuracy: 0.8018
                 95% CI : (0.7898, 0.8133)
   No Information Rate: 0.784
   P-Value [Acc > NIR] : 0.00183
                  Kappa: 0.3701
Mcnemar's Test P-Value : 2.481e-10
            Sensitivity: 0.9005
            Specificity: 0.4434
         Pos Pred Value: 0.8545
         Neg Pred Value: 0.5512
             Prevalence: 0.7840
         Detection Rate: 0.7060
  Detection Prevalence: 0.8262
     Balanced Accuracy: 0.6720
```

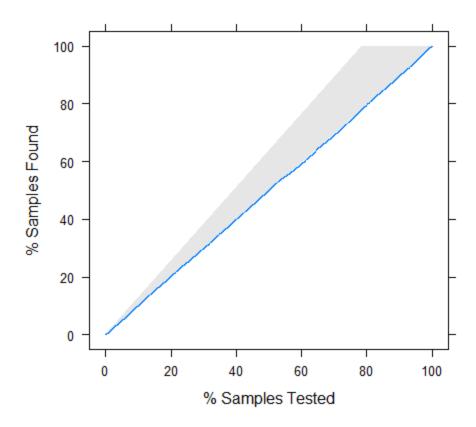
'Positive' Class: 0

ROC curve



... Precision/Recall graphs ...





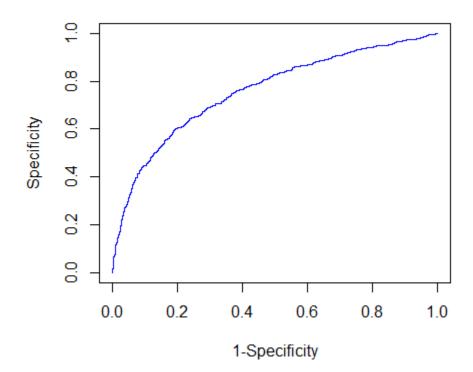
Neural Networks

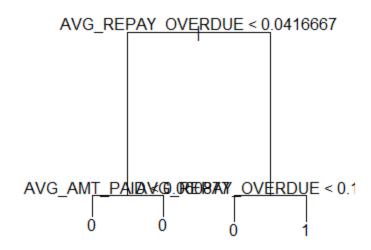
- 1. To go ahead with Neural Networks model we first cleaned the data. Removed the outliers. Plotted various graphs such as Histograms, Box plot to maintain the integrity of data.
- 2. Next we sampled the data into Test and Train to model the train data and then evaluate our model on based of Test data set.
- 3. In Neural Networks we did modeling with variables : LIMIT_BAL+SEX+MARRIAGE+AGE+AVG_REPAY_OVERDUE+AVG_AMT_PAID+AVG_BIL L_REPAY
- 4. Graphs were plotted to evaluate the performance of the models

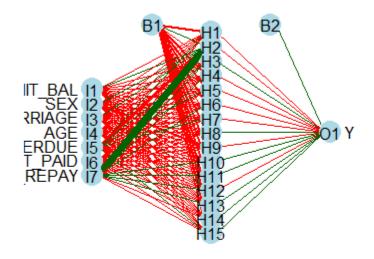
```
library(NeuralNetTools)
maxs \leftarrow apply(x, 2, max)
mins \leftarrow apply(x, 2, min)
scaled <- as.data.frame(scale(x, center = mins, scale = maxs - mins))</pre>
index \leftarrow sample(1:nrow(x), round(0.90*nrow(x)))
train_ <- scaled[index,]</pre>
test_ <- scaled[-index,]</pre>
test_$Y <- as.factor(test_$Y)</pre>
train_$Y <- as.factor(train_$Y)</pre>
# library(neuralnet)
# n <- names(train_)</pre>
# f <- as.formula(paste("Y ~", paste(n[!n %in% "Y"], collapse = " + ")))
# nn <- neuralnet(f,data=train_,hidden=c(4,3))</pre>
#Build the model
fitnn <- nnet(Y~LIMIT_BAL+SEX+MARRIAGE+AGE+AVG_REPAY_OVERDUE+AVG_AMT_PAID+AVG_BILL_REPAY,
              data=train_{,} size=15, hess = T, dk=5e-4, maxit = 200)
summary(fitnn)
#Predict for test data
pred = predict(fitnn, newdata=test_, type="class")
str(test_)
#Confucion Matrix
> # ROC area under the curve
> pred = predict(fitnn, newdata=test_, type="raw")
> prediction <- prediction(pred,test_$Y)
> auc.tmp <- performance(prediction, "auc")
> auc <- as.numeric(auc.tmp@y.values)</pre>
> print(auc)
 [1] 0.7595364
```

```
> library(e1071)
> confusionMatrix(pred,test_$Y)
Confusion Matrix and Statistics
          Reference
         on 0 1
0 2170 397
1 170 263
Prediction
               Accuracy: 0.811
                 95% CI: (0.7965, 0.8249)
    No Information Rate : 0.78
    P-Value [Acc > NIR] : 1.691e-05
                  Kappa : 0.3717
 Mcnemar's Test P-Value : < 2.2e-16
            Sensitivity: 0.9274
            Specificity: 0.3985
         Pos Pred Value: 0.8453
         Neg Pred Value: 0.6074
             Prevalence: 0.7800
         Detection Rate: 0.7233
   Detection Prevalence: 0.8557
      Balanced Accuracy: 0.6629
       'Positive' Class : 0
```

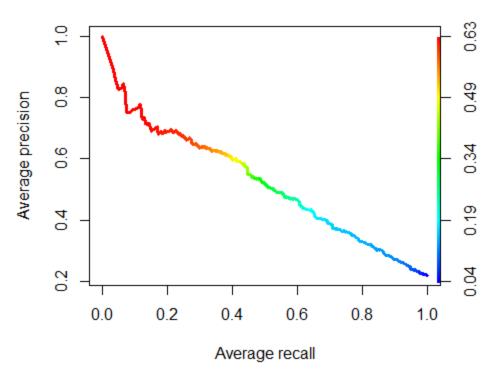
ROC curve

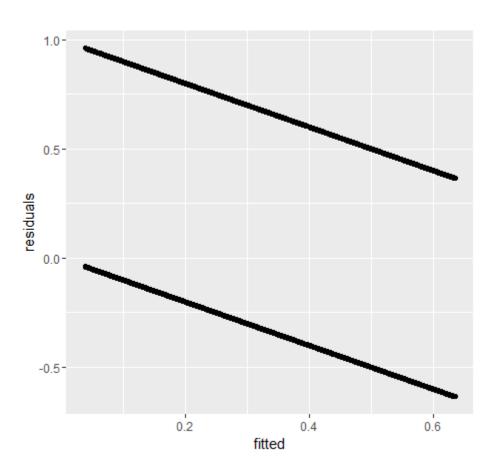


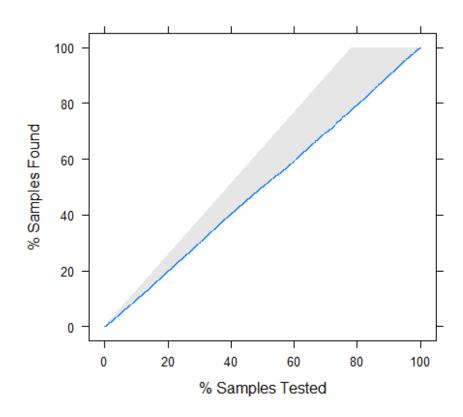




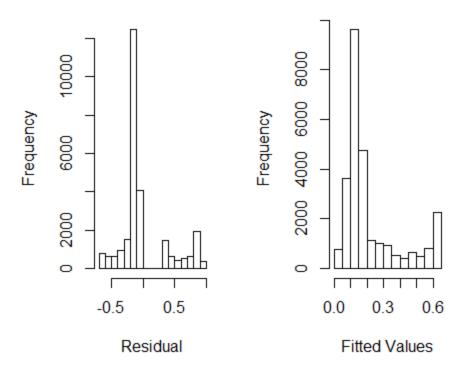
... Precision/Recall graphs ...







Histogram of Resdiual: Histogram of Fitted Valu



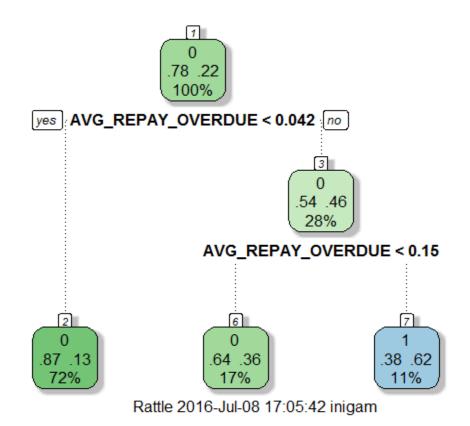
Classification Tree

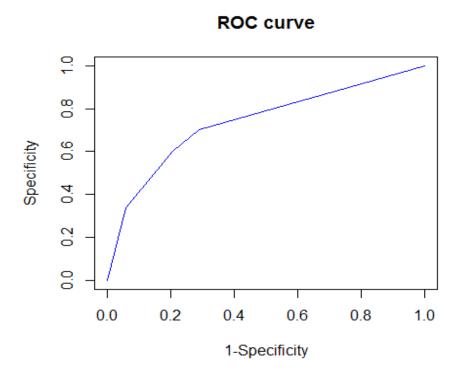
- 1. To go ahead with classification model we first cleaned the data. Removed the outliers. Plotted various graphs such as Histograms, Box plot to maintain the integrity of data.
- 2. Next we sampled the data into Test and Train to model the train data and then evaluate our model on based of Test data set.
- 3. In classification we did modeling with variables all variables
- 4. Graphs were plotted to evaluate the performance of the models

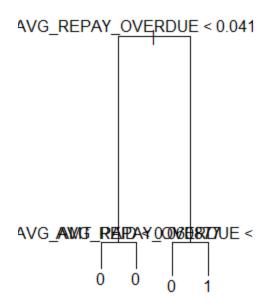
```
library(tree)
tree.train = tree(Y~.,data=train_)
summary(tree.train)

#Display the tree structure and node labels
plot(tree.train)
text(tree.train, pretty =0) #Pretty=0 includes the category names
```

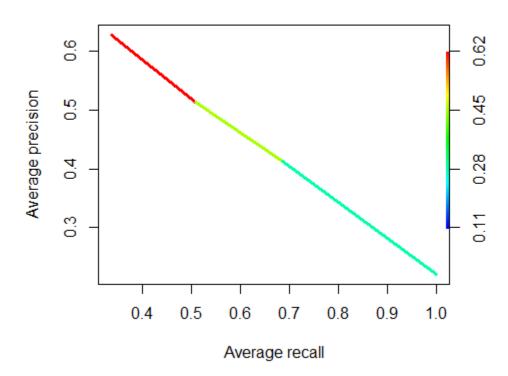
```
> # ROC area under the curve
> tree.pred = predict(tree.train, test_, type ="vector")
> prediction <- prediction(tree.pred[,2],test_$Y)</pre>
> auc.tmp <- performance(prediction, "auc")</pre>
> auc <- as.numeric(auc.tmp@y.values)</pre>
> print(auc)
[1] 0.7388666
> confusionmatrix(tree.pred,test_$Y)
Confusion Matrix and Statistics
          Reference
Prediction 0 1
         0 2208 437
         1 132 223
               Accuracy: 0.8103
                 95% CI: (0.7958, 0.8242)
    No Information Rate: 0.78
    P-Value [Acc > NIR] : 2.511e-05
                  Kappa : 0.3374
 Mcnemar's Test P-Value : < 2.2e-16
            Sensitivity: 0.9436
            Specificity: 0.3379
         Pos Pred Value : 0.8348
         Neg Pred Value: 0.6282
             Prevalence: 0.7800
         Detection Rate: 0.7360
   Detection Prevalence: 0.8817
      Balanced Accuracy: 0.6407
       'Positive' Class: 0
tree.pred 0 1
      0 2208 437
       1 132 223
Classification tree:
tree(formula = Y \sim ., data = train_)
Variables actually used in tree construction:
[1] "AVG_REPAY_OVERDUE" "AVG_AMT_PAID"
Number of terminal nodes: 4
Residual mean deviance: 0.902 = 24350 / 27000
Misclassification error rate: 0.1939 = 5236 / 27000
> |
```

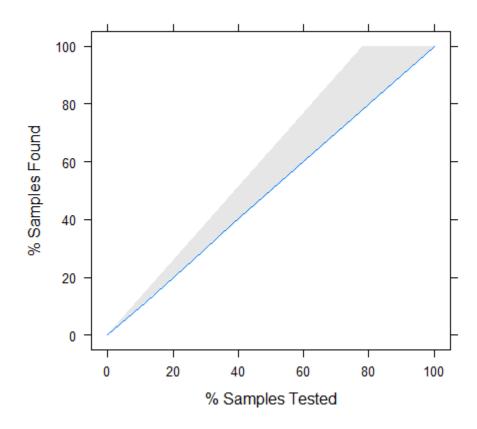






... Precision/Recall graphs ...





Discussion

- ✓ With the results derived above we can conclude that Neural Network
 has best Performance metrics, the sensitivity and specificity hold best
 results along with high range of accuracy
- ✓ Also, the area under ROC curve and lift curve analysis supports neural network model the most accurate model on given dataset.
- ✓ In neural network, the prediction of default payment for next month was predicted agrees with the default payment for next month provided with dataset, can be proved with confusion matrix.

Problem2- Advertisements

4 GOAL

- 1. The dataset contains a set of advertisements found on Internet. These advertisements are mostly described in image geometries as well as phrases contained in the URL, the image's URL, alt text, the anchor text, and places near the anchor text. The dataset is a two-class data set. In other word, its class attribute is Boolean, which is either "ad" or "nonad", corresponds to "It is an advertisement" and "It is not an advertisement". There are 3279 instances, where 2821 instances are "nonads" and 458 instances are "ads". There are 1558 attributes, where 3 attributes are continuous and others are binary.
- 2. Our Goal is to clean the data, process it to make sure the outliers are removed and then apply logistic regression, neural network and classification tree to predict and classify the model
- 3. We have to build different models to evaluate different trends and compare the sensitivity, Specification and plot confusion matrix, ROC Curve and Lift Charts for all the classification method.

Data Cleaning

- 1. It was observed 28% of continuous data was missing which are dimensions of advertisements and area ratio. There are various results to fill the missing data.
 - ✓ Replace by mean
 - ✓ Replace by median
 - ✓ Replace with the previous value using locf

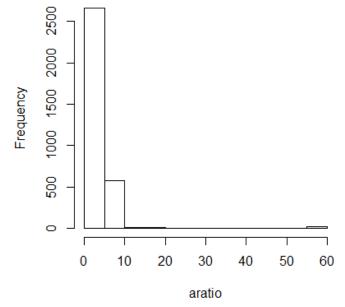
Approach:- Replacing the values with locf will not make sense as previous values has no relation with each other as each ad is different. Replacing the values with mean or median works best so as to make sure the data set is concentrated towards a common area. We choose mean to replace it and plotted histogram.

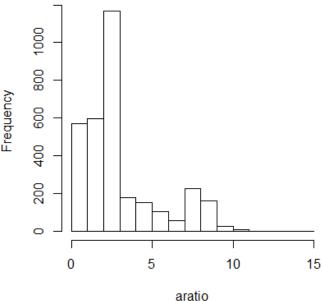
2. Removing the outliers: - Plotted histogram as shown below to see the range to data where they lie. We choose frequency of aspect ratio as a base as it is calculated using height and width. From the histogram as well as from power BI, it's clear that the data is more frequently occurring with aspect ratio of less than 12. So we remove the data which had aspect ratio greater than 12.

```
#Replace ? with NA
addata$height[addata$height == " ?"] <- NA
addata$width[addata$width == " ?"] <- NA
addata$aratio[addata$aratio == "
#Convert charcters to numerics to calculate mean
addata$ad.NonAd <- as.factor(addata$ad.NonAd)</pre>
addata$height <- as.numeric(addata$height)
addata$width <- as.numeric(addata$width)</pre>
addata$aratio <- as.numeric(addata$aratio)</pre>
#calculating mean
meanHeight <- mean(addata$height, na.rm=TRUE)</pre>
meanWidth <- mean(addata$width, na.rm=TRUE)</pre>
#replacing missing data with mean
addata$height[is.na(addata$height)] = meanHeight
addata$width[is.na(addata$width)] = meanWidth
addata$aratio[is.na(addata$aratio)] = meanWidth/meanHeight
```

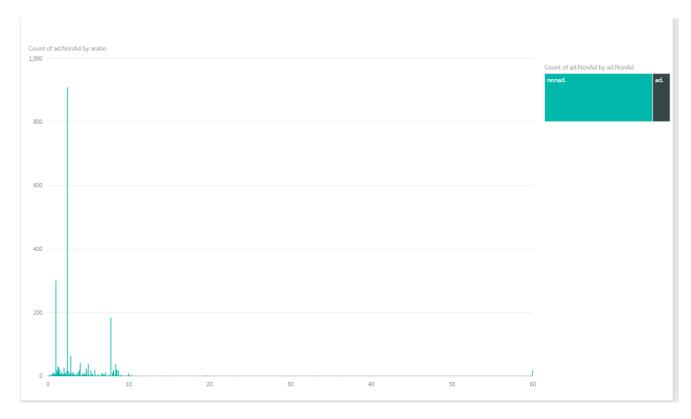
Histogram with outliers

Histogram without outliers

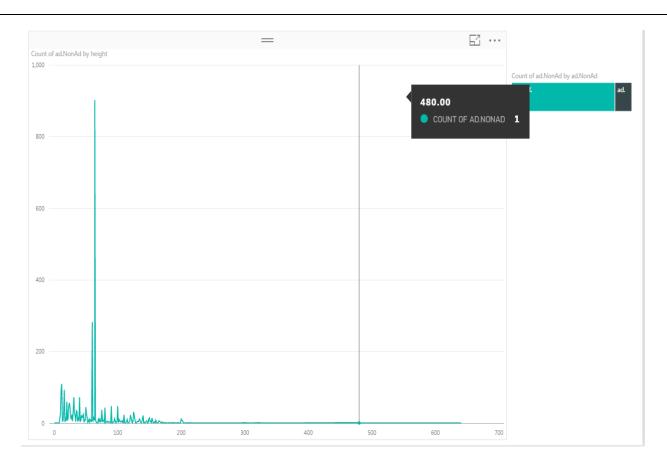




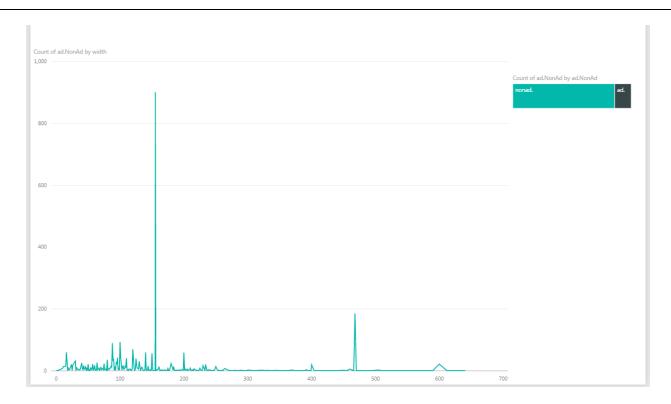
♣Power PI and Inferences



✓ The above graph shows count of Non-Ads/Ads by Area Ratio. We can depict that Non add whose total count is 2820 and major of the non-ads have high are ratio less than 12 and all of the Ads lie in area ratio below 12.



✓ We can conclude that almost all non-ads also lie with under height of 200 and Ads have kind of constant height throughout so we can conclude that almost all the Ads have same height.



- ✓ We can observe that all Ads have variety of distribution when it comes to height but in case of widths it has distribution throughout.
- ✓ Using tool Power BI we can conclude the with height, width and area ratio which are continuous data set in our problem, we can conclude that all the Ads have area ratio less than 12. So we can consider aspect ratio to predict whether an Advertisement is Ad or Non-Ad

Logistic Regression

- 1. To go ahead with logistic regression model we first cleaned the data. Next we sampled the data into Test and Train to model the train data and then evaluate our model on based of Test data set. We have taken 75% train data but according to a research study, its better to take 90% train data and 10% test data as we don't want a ad to be classified as non-ad as it is very signification according to business scenario.
- 2. In Logistic regression we did GLM modeling with variables url.ad+alt.click+url.images.home+height+ancurl.com+ ancurl.exe+url.ads+alt.net+width+ancurl.click+ancurl.http.www+aratio
- 3. Threshold of 0.5 probability was set to classify data into add or non-ad and predictions were made as per the analysis.
- 4. Graphs were plotted to evaluate the performance of the models

```
> confusionMatrix(pred,test$ad.NonAd)
Confusion Matrix and Statistics
          Reference
Prediction ad. nonad.
    ad.
          107
   nonad. 15
                  682
              Accuracy: 0.9716749
                 95% CI: (0.9578003, 0.9819614)
   No Information Rate: 0.8497537
   P-Value [Acc > NIR] : < 0.00000000000000022
                  Kappa: 0.886388
Mcnemar's Test P-Value: 0.2109029
            Sensitivity: 0.8770492
            Specificity: 0.9884058
         Pos Pred Value : 0.9304348
         Neg Pred Value : 0.9784792
             Prevalence : 0.1502463
         Detection Rate: 0.1317734
  Detection Prevalence: 0.1416256
      Balanced Accuracy: 0.9327275
      'Positive' Class : ad.
```


AUC:-


From a random classifier you can expect as many true positives as false positives. That's the dashed line on the plot. AUC score for the case is 0.5. A score for a perfect classifier would be 1.

```
> # ROC area under the curve
> prediction <- prediction(prob,test$ad.NonAd)
> auc.tmp <- performance(prediction,"auc")
> auc <- as.numeric(auc.tmp@y.values)
> print(auc)
[1] 0.9760572583
```

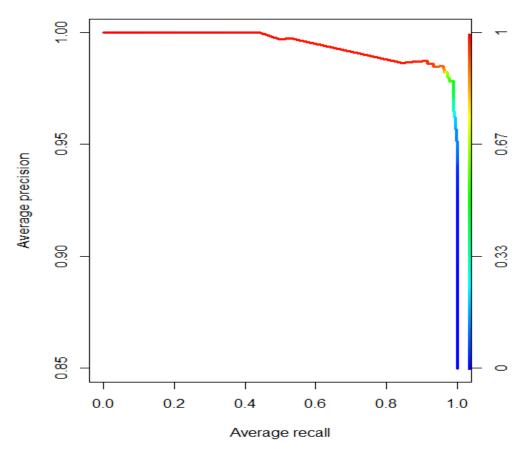
ROC:-

In statistics, a receiver operating characteristic (ROC), or ROC curve, is a graphical plot that illustrates the performance of a binary classifier system as its discrimination threshold is varied. The curve is created by plotting the true positive rate (TPR) against the false positive rate (FPR) at various threshold settings. The true-positive rate is also known as sensitivity, or recall in machine learning. The false-positive rate is also known as the fall-out and can be calculated as (1 - specificity). An ideal ROC is the one which is closer to top left corner as below.









4Classification Tree

- 1. To go ahead with classification model we first cleaned the data. Removed the outliers. Plotted various graphs such as Histograms, Box plot to maintain the integrity of data.
- 2. Next we sampled the data into Test and Train to model the train data and then evaluate our model on based of Test data set.
- 3. In classification we did modeling with variables all variables
- 4. Graphs were plotted to evaluate the performance of the models

```
#Use variables to fit a classification tree
library(tree)
tree.train = tree(ad.NonAd~.,data=train)
summary(tree.train)
#Display the tree structure and node labels
plot(tree.train)
text(tree.train, pretty =0) #Pretty=0 includes the category names
#FancyRPlot
library(rpart)
library(rattle)
tree <- rpart(ad.NonAd~.,data=train,method="class")</pre>
fancyRpartPlot(tree)
tree.pred = predict(tree.train, test, type ="class")
table(tree.pred, test$ad.NonAd)
> confusionMatrix(tree.pred,test$ad.NonAd)
Confusion Matrix and Statistics
           Reference
Prediction ad. nonad.
           106
    ad.
                    16
    nonad. 16
                   674
                Accuracy : 0.9605911
                  95% CI: (0.9448196, 0.972891)
    No Information Rate: 0.8497537
    Kappa: 0.8456641
Mcnemar's Test P-Value : 1
             Sensitivity: 0.8688525
             Specificity: 0.9768116
          Pos Pred Value : 0.8688525
          Neg Pred Value : 0.9768116
              Prevalence : 0.1502463
          Detection Rate: 0.1305419
   Detection Prevalence: 0.1502463
      Balanced Accuracy: 0.9228320
       'Positive' Class : ad.
> summary(tree.train)
classification tree:
tree(formula = ad.NonAd ~ ., data = train)
Variables actually used in tree construction:
 [1] "width" "ancurl.o
[7] "url.images.home" "alt.net'
                     "ancurl.com"
                                     "url.ads"
                                                      "ancurl.click"
                                                                      "url.ad"
                                                                                      "alt.click"
                                     "aratio"
                                                      "height"
Number of terminal nodes: 12
Residual mean deviance: 0.2214256 = 536.7356 / 2424
Misclassification error rate: 0.02873563 = 70 / 2436
```

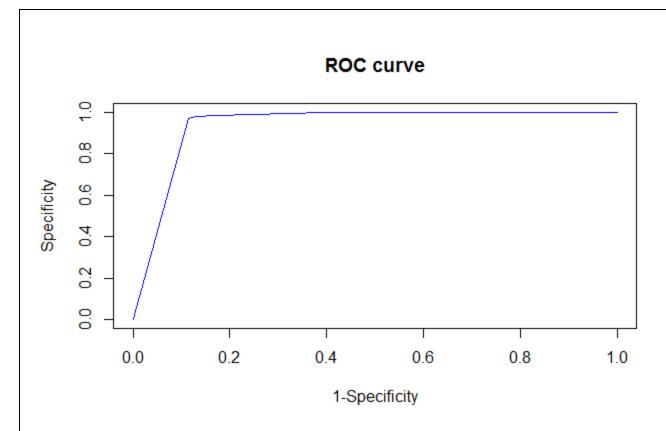
AUC:-

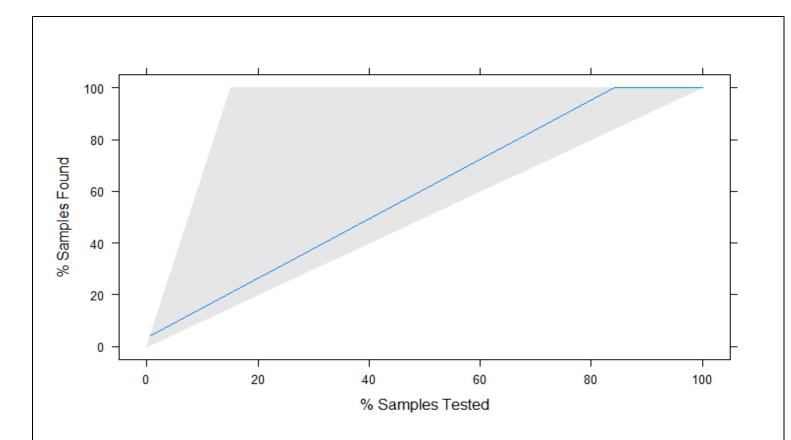
From a random classifier you can expect as many true positives as false positives. That's the dashed line on the plot. AUC score for the case is 0.5. A score for a perfect classifier would be 1.

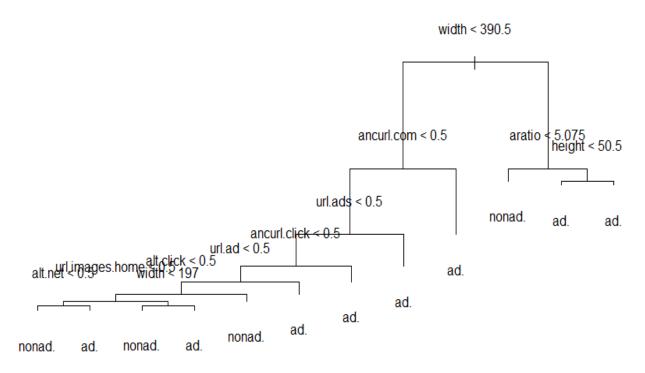
```
> prediction <- prediction(tree.pred[,2],test$ad.NonAd)
> auc.tmp <- performance(prediction,"auc")
> auc <- as.numeric(auc.tmp@y.values)
> print(auc)
[1] 0.9377464956
```

ROC:-

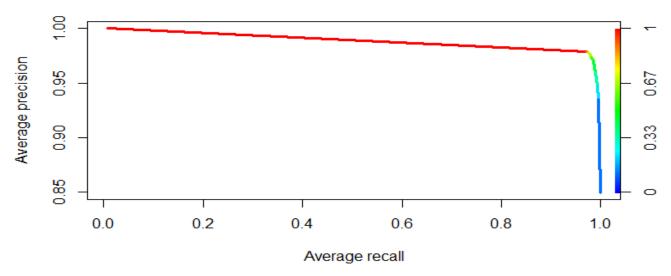
In statistics, a receiver operating characteristic (ROC), or ROC curve, is a graphical plot that illustrates the performance of a binary classifier system as its discrimination threshold is varied. The curve is created by plotting the true positive rate (TPR) against the false positive rate (FPR) at various threshold settings. The true-positive rate is also known as sensitivity, or recall in machine learning. The false-positive rate is also known as the fall-out and can be calculated as (1 - specificity). An ideal ROC is the one which is closer to top left corner as below.














Neural Networks

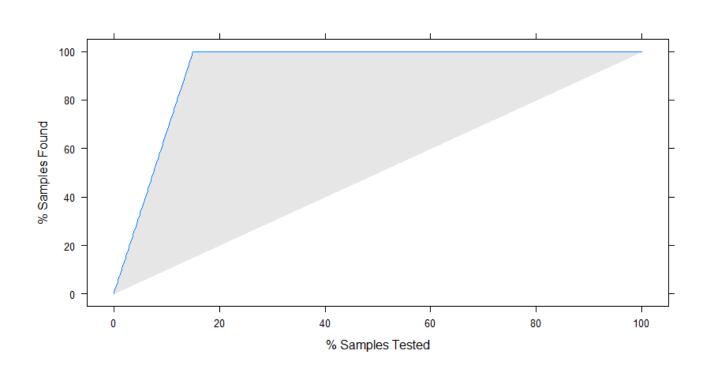
- 1. To go ahead with Neural Networks model we first cleaned the data. Removed the outliers. Plotted various graphs such as Histograms, Box plot to maintain the integrity of data.
- 2. In Neural Networks we did modeling with variables url.ad+alt.click+url.images.home+ancurl.com+height+ancurl.exe+ url.ads+alt.net+width+ancurl.click+ancurl.http.www
- 3. Graphs were plotted to evaluate the performance of the models


```
#Neural Network
          library(nnet)
          library(NeuralNetTools)
         #Build the model
         fitnn <- nnet(ad.NonAd~url.ad+alt.click+url.images.home+ancurl.com+height+ancurl.exe+
                        url.ads+alt.net+width+ancurl.click+ancurl.http.www,data=train, size=4,
                       hess = T, dk=15e-4, maxit = 200)
         summary(fitnn)
         #Predict for test data
         pred = predict(fitnn, newdata=test, type="class")
          #Confusion Matrix
          library(caret)
          library(e1071)
         confusionMatrix(pred,test$ad.NonAd)
> summary(fitnn)
a 11-4-1 network with 53 weights
options were - entropy fitting
                                         i5->h1 i6->h1 i7->h1
                                                                 i8->h1
  b->h1
        i1->h1 i2->h1
                        i3->h1
                                 i4->h1
                                                                         i9->h1 i10->h1 i11->h1
                                   0.23
                                         -5.11
                                                  -0.48
                                                          -0.62
                                                                  -0.68
                                                                                   -0.33
  -2.27
         -2.44
                 -0.62
                           0.42
                                                                          -4.07
                                                                                           -0.62
                                i4->h2
                                                 i6->h2
                                                         i7->h2
                                                                         i9->h2 i10->h2 i11->h2
  b->h2 i1->h2
                 i2->h2
                         i3->h2
                                         i5->h2
                                                                 i8->h2
         -8.57
                 -3.82
                         -3.28
                                -50.34
                                                  -5.47
                                                         -23.38
                                                                   6.46
                                                                          -0.02 -40.14
                                                                                         -21.67
  4.32
                                           0.03
                                                 i6->h3 i7->h3
                                                                         i9->h3 i10->h3 i11->h3
  b->h3
        i1->h3
                i2->h3
                        i3->h3
                                 i4->h3
                                         i5->h3
                                                                 i8->h3
  -2.04
          0.01
                   0.66
                         -0.42
                                  -0.09
                                         -1.97
                                                  -0.65
                                                          -0.56
                                                                  -0.11
                                                                          -2.38
                                                                                   -0.66
                                         i5->h4 i6->h4 i7->h4
  b->h4 i1->h4 i2->h4
                        i3->h4
                                i4->h4
                                                                 i8->h4 i9->h4 i10->h4 i11->h4
         11.62
                 -0.23
                         -0.01
                                   0.35
                                         -1.06
                                                   0.48
                                                         -0.25
                                                                  -0.37
                                                                          -2.28
                                                                                   0.41
  b->0 h1->0 h2->0 h3->0 h4->0
 -2.90 -1.70 7.03 -2.24 -11.55
```

AUC:-

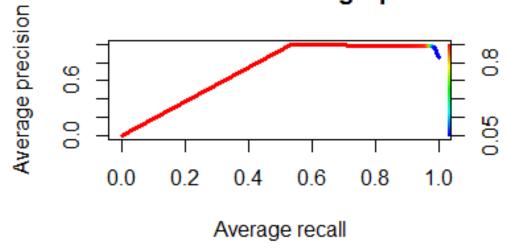
From a random classifier you can expect as many true positives as false positives. That's the dashed line on the plot. AUC score for the case is 0.5. A score for a perfect classifier would be 1.

```
> # ROC area under the curve
 > pred = predict(fitnn, newdata=test, type="raw")
 > prediction <- prediction(pred,test$ad.NonAd)</pre>
 > auc.tmp <- performance(prediction, "auc")</pre>
 > auc <- as.numeric(auc.tmp@y.values)</pre>
 > print(auc)
 [1] 0.9602934189
> contusionMatrix(pred,test$ad.NonAd)
Confusion Matrix and Statistics
          Reference
Prediction ad. nonad.
          109
                   17
    ad.
    nonad.
           13
                  673
               Accuracy: 0.9630542
                 95% CI: (0.9476756, 0.9749366)
    No Information Rate: 0.8497537
    P-Value [Acc > NIR] : < 0.00000000000000022
                  Kappa: 0.8572366
 Mcnemar's Test P-Value: 0.5838824
            Sensitivity: 0.8934426
            Specificity: 0.9753623
         Pos Pred Value : 0.8650794
         Neg Pred Value : 0.9810496
             Prevalence : 0.1502463
         Detection Rate: 0.1342365
   Detection Prevalence: 0.1551724
      Balanced Accuracy: 0.9344025
       'Positive' Class: ad.
```





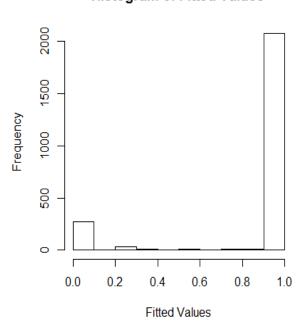
... Precision/Recall graphs ...




Histogram of Resdiuals

Leadneuch -1.0 -0.5 0.0 0.5 1.0 Residual

Histogram of Fitted Values





Discussion

- ✓ With the results derived above we can conclude that Classification Tree
 has good Performance metrics, the sensitivity and specificity hold best
 results along with high range of accuracy
- ✓ Also, the area under ROC curve and lift curve analysis supports Classification Tree the most accurate model on given dataset.
- ✓ In Classification Tree, the prediction of ad and non-ad classification is given in the confusion matrix.

Problem3- Wind Forecasting

4 GOAL

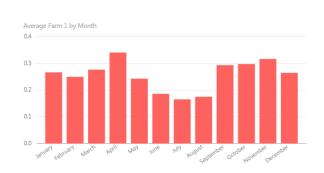
- 1. The dataset contains wind power forecast value for seven different warms over 4.5 years of date.
- 2. Our Goal is to clean the data, process it to make sure the data is clean and without outliners. There is test and train data based on date filter. We have to make model using linear regression, neural network and Regression tree to forecast the data in Test data for missing data.
- 3. We have to build different models to evaluate different trends and compare the performance metric for all the models.

Data Cleaning

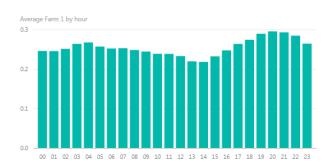
- 1. It was observed data had date with time stamp and there were multiple entries with same timestamp so we took mean of all the entries for Single time stamp.
- 2. Histogram was plotted to see outliers and using Cook's Method we removed outliers.
- 3. The complicated part of data cleaning was we have to separate and refine the data with its time stamp for all the seven farm files and then later we had to merge all the Wp(wind power) into one single file with Timestamps.

```
for(i in 1:7){
 assign(paste0("wf"), read.csv(paste("windforecasts_wf",
                                     i, ".csv", sep=""),colClasses = c("character", "numeric",
 wf\$date <- as.POSIXct(wf\$date,format = "%Y%m%d%H")
 wf$actualDate <- wf$date + (wf$hors * 3600)
 wf <- wf[,c(3:7)]
 meanU <- mean(wf$u, na.rm=TRUE)
 meanV <- mean(wf$v, na.rm=TRUE)
 meanWS <- mean(wf$ws, na.rm=TRUE)
 meanWD <- mean(wf$wd. na.rm=TRUE)
 wf$u[is.na(wf$u)] = meanU
 wf v[is.na(wf v)] = meanv
 wf$ws[is.na(wf$ws)] = meanWS
 wf$wd[is.na(wf$wd)] = meanWD
 wf <- aggregate(x = wf[c("u","v","ws","wd")],</pre>
                 FUN = mean.
                 by = list(Date = wf$actualDate))
 wf$hour <- as.factor(format(wf$Date, "%H"))</pre>
 wf$OnlyDate <- as.Date(format(wf$Date))</pre>
 assign(paste0("wf", i), wf)
```

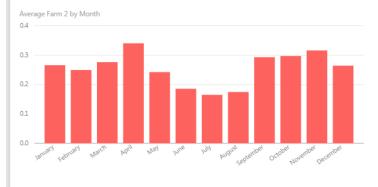
♣Power B I and Inferences





FARM 1
















FARM 3




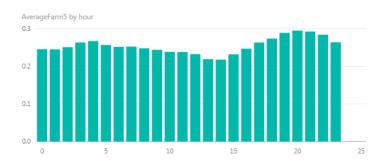






FARM 4












FARM 5











- ✓ As per the above graphs we can see a pattern in the Average wind power to month and Hours for all the 7 farms.
- ✓ The Wind power is maximum in April month and in month of May-September the wind power decreases.
- ✓ The wind power increases from September to November.
- ✓ On per hour basis the wind power is observed to be maximum in evening Hours and decreases around afternoon 1pm to 3pm.
- ✓ The wind speed and wind direction and zonal and meridional wind
 components also have a =n impact on the output which is wind power and
 temperature, hour, month such factors impact the predictors u,v,ws,wd

Linear Regression

- 1. To go ahead with linear regression model we first cleaned the data and removed outliers with cook's method. Next we sampled the data into Test and Train to model the train data and then evaluate our model on based of Test data set.
- 2. In linear regression we did modeling for all 7 farms and the models for all the farms were different.
- 3. Feature Selection was done all the farms and model was applied based on the different predictors.
- 4. Performance metrics and accuracy was calculated for all seven models.

```
[]m.fit1=]m(wp1~.-Date-OnlyDate-month-year,data=model1)
summary(lm.fit1)
pred = predict(lm.fit1, test1)
accuracy(pred, test1$wp1)
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -7.603e-02 5.239e-03 -14.513 < 2e-16 ***
            -5.711e-03 2.030e-04 -28.141 < 2e-16 ***
            -4.473e-03 8.016e-04 -5.580 2.45e-08 ***
u
             8.367e-03 4.974e-04 16.822 < 2e-16 ***
             1.026e-01 8.835e-04 116.081 < 2e-16 ***
WS
wd
            -1.449e-04 2.539e-05 -5.706 1.18e-08 ***
Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' ' 1
Residual standard error: 0.1552 on 13127 degrees of freedom
Multiple R-squared: 0.5746, Adjusted R-squared: 0.5744
F-statistic: 3546 on 5 and 13127 DF, p-value: < 2.2e-16
> pred = predict(lm.fit1, test1)
> accuracy(pred, test1$wp1)
                        RMSE
                                  MAE MPE MAPE
Test set 0.0356211 0.1773012 0.1339909 NaN Inf
```

```
lm.fit2=lm(wp2~.-Date-OnlyDate-month-u, data=model2)
summary(1m.fit2)
library(forecast)
pred = predict(lm.fit2, test2)
accuracy(pred, test2$wp2)
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -1.264e-01 7.130e-03 -17.73 <2e-16 ***
                                            <2e-16 ***
           -7.131e-03 2.207e-04 -32.31
hour
                                          <2e-16 ***
            1.086e-02 5.005e-04
                                  21.70
            1.138e-01 8.387e-04
5.754e-04 1.659e-05
                                           <2e-16 ***
                                  135.68
W.S
                                           <2e-16 ***
wd
                                   34.69
            -4.491e-02 3.232e-03 -13.89 <2e-16 ***
year
Signif. codes: 0 '*** 0.001 '** 0.01 '* 0.05 '.' 0.1 ' ' 1
Residual standard error: 0.17 on 13064 degrees of freedom
Multiple R-squared: 0.5999, Adjusted R-squared: 0.5997
F-statistic: 3918 on 5 and 13064 DF, p-value: < 2.2e-16
> library(forecast)
> pred = predict(lm.fit2, test2)
> accuracy(pred, test2$wp2)
                 MF
                         RMSE
                                    MAE MPE MAPE
Test set 0.06563687 0.1819857 0.1387788 NaN Inf
```

```
lm.fit3=lm(wp3~.-Date-OnlyDate-year-month-wd, data=model3)
summary(1m.fit3)
library(forecast)
pred = predict(1m.fit3, test3)
accuracy(pred, test3$wp3)
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
                                          <2e-16 ***
(Intercept) -0.1357297 0.0040748 -33.31
                                           <2e-16 ***
           -0.0070628 0.0002145 -32.92
hour
            -0.0058941 0.0004309 -13.68
                                           <2e-16 ***
u
                                           <2e-16 ***
            0.0111887 0.0004329 25.84
                                          <2e-16 ***
            0.1205050 0.0006978 172.70
WS
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' '1
Residual standard error: 0.1677 on 13100 degrees of freedom
Multiple R-squared: 0.7129,
                               Adjusted R-squared: 0.7128
F-statistic: 8130 on 4 and 13100 DF, p-value: < 2.2e-16
> library(forecast)
> pred = predict(lm.fit3, test3)
> accuracy(pred, test3$wp3)
                ME
                        RMSE
                                   MAE MPE MAPE
Test set 0.02410656 0.1862717 0.1407924 NaN Inf
```

```
lm.fit4=lm(wp4~.-Date-OnlyDate-year-u-month, data=model4)
summary(lm.fit4)
library(forecast)
pred = predict(lm.fit4, test4)
accuracy(pred, test4$wp4)
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -1.714e-01 4.935e-03 -34.728 <2e-16 ***
            -7.181e-03 2.016e-04 -35.628
1.233e-02 3.937e-04 31.327
                                             <2e-16 ***
hour
                                             <2e-16 ***
                                            <2e-16 ***
             1.143e-01 6.661e-04 171.630
W.S
                                             <2e-16 ***
wd
             1.531e-04 1.586e-05
                                    9.652
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
Residual standard error: 0.1563 on 13125 degrees of freedom
Multiple R-squared: 0.703,
                                Adjusted R-squared: 0.703
F-statistic: 7768 on 4 and 13125 DF, p-value: < 2.2e-16
> library(forecast)
> pred = predict(lm.fit4, test4)
> accuracy(pred, test4$wp4)
                 MF
                         RMSE
                                     MAE MPE MAPE
Test set 0.02954673 0.1707248 0.1306941 NaN Inf
Wind Farm 5
lm.fit5=lm(wp5\sim.-Date-OnlyDate-month+I(u^2)+I(wd^2), data=model5)
summary(1m.fit5)
library(forecast)
pred = predict(lm.fit5, test5)
accuracy(pred, test5$wp5)
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -3.168e-01 9.746e-03 -32.508 < 2e-16 ***
            -3.489e-03 2.006e-04 -17.396 < 2e-16 ***
hour
            -1.255e-02 7.848e-04 -15.995 < 2e-16 ***
1.607e-02 7.015e-04 22.903 < 2e-16 ***
u
             1.076e-01 9.915e-04 108.566 < 2e-16 ***
W.S
                                   6.611 3.97e-11 ***
wd
             6.542e-04 9.895e-05
            6.499e-02 2.929e-03 22.187 < 2e-16 ***
year
             2.110e-03
                        1.375e-04 15.342 < 2e-16 ***
I(u^2)
            -2.599e-06 2.842e-07 -9.144 < 2e-16 ***
I(wd^2)
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' '1
Residual standard error: 0.1566 on 13066 degrees of freedom
Multiple R-squared: 0.6835, Adjusted R-squared: 0.6833
F-statistic: 3527 on 8 and 13066 DF, p-value: < 2.2e-16
> library(forecast)
> pred = predict(lm.fit5, test5)
> accuracy(pred, test5$wp5)
                          RMSE
                                     MAE MPE MAPE
                  MF
Test set -0.05783443 0.1918051 0.1555689 NaN Inf
```

```
lm.fit6=lm(wp6~.-Date-OnlyDate-year-u, data=model6)
summary(lm.fit6)
library(forecast)
pred = predict(lm.fit6, test6)
accuracy(pred, test6$wp6)
Coefficients:
               Estimate Std. Error t value Pr(>|t|)
(Intercept) -1.898e-01 5.455e-03 -34.797 < 2e-16 ***
             -4.279e-03 1.871e-04 -22.874 < 2e-16 ***
             7.734e-03 3.159e-04 24.482 < 2e-16 ***
9.892e-02 5.608e-04 176.412 < 2e-16 ***
1.009e-04 1.462e-05 6.901 5.42e-12 ***
V
WS
wd
              1.029e-03 3.907e-04 2.633 0.00847 **
month
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
Residual standard error: 0.1457 on 13095 degrees of freedom
Multiple R-squared: 0.7125, Adjusted R-squared: 0.7124
F-statistic: 6490 on 5 and 13095 DF, p-value: < 2.2e-16
> library(forecast)
> pred = predict(lm.fit6, test6)
> accuracy(pred, test6$wp6)
                  ME
                           RMSE
Test set 0.04763217 0.1678523 0.1241318 NaN Inf
```

```
lm.fit7 = lm(wp7 \sim .- Date-OnlyDate-u-year-month, \ data = model7)
 summary(lm.fit7)
 library(forecast)
 pred = predict(lm.fit7, test7)
 accuracy(pred, test7$wp7)
Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -2.559e-01 4.963e-03 -51.56 <2e-16 ***
                                          <2e-16 ***
           -2.944e-03 1.948e-04 -15.11
hour
            1.159e-02 2.891e-04 40.10 <2e-16 ***
            9.700e-02 5.048e-04 192.17 <2e-16 ***
WS
            2.647e-04 1.512e-05
                                  17.50 <2e-16 ***
wd
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
Residual standard error: 0.1521 on 13092 degrees of freedom
Multiple R-squared: 0.7437, Adjusted R-squared: 0.7436
F-statistic: 9497 on 4 and 13092 DF, p-value: < 2.2e-16
> library(forecast)
> pred = predict(lm.fit7, test7)
> accuracy(pred, test7$wp7)
                       RMSE
                                  MAE MPE MAPE
                ME
Test set 0.01396243 0.164149 0.1257243 NaN Inf
> |
```

♣Neural Network

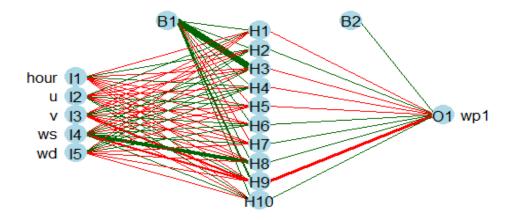
- 1. To go ahead with neural network model we first cleaned the data and removed outliers with cook's method.
- 2. In Neural network we did modeling for all 7 farms and the models for all the farms were different.
- 3. Feature Selection was done all the farms and model was applied based on the different predictors.
- 4. Performance metrics RMSE was calculated for all seven models.

Wind Farm 1

```
library(forecast)
library(ROCR)
fml<- as.formula("wpl~.-Date-OnlyDate-month-year");
res <- nnet(fml, data=model1,size=10, linout=TRUE, skip=TRUE, MaxNWts=10000, trace=FALSE, maxit=100)
summary(res)

#plot
plotnet(res,circle_col = "lightblue", circle_cex = 3,pos_col = "darkgreen", neg_col = "red")

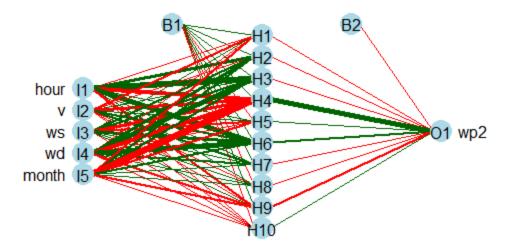
#predict
pred<-predict(res, newdata=test1,type="raw")
pred <- round(pred, digits = 3)
rmse <- sqrt(mean((pred- test1$wp1)^2))
mape <- mean(abs((test1$wp1 - pred)/test1$wp1))*100</pre>
```

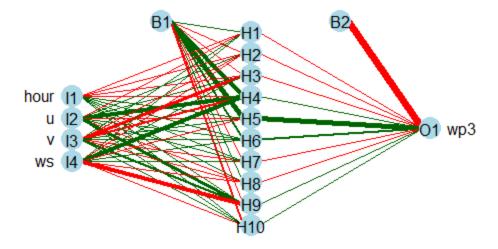


```
fml<- as.formula("wp2~.-Date-OnlyDate-year-u");
res <- nnet(fml, data=model2,size=10, linout=TRUE, skip=TRUE, MaxNWts=10000, trace=FALSE, maxit=100)
summary(res)

#plot
plotnet(res,circle_col = "lightblue", circle_cex = 3,pos_col = "darkgreen", neg_col = "red")

#predict
pred<-predict(res, newdata=test2,type="raw")
pred <- round(pred, digits = 3)
rmse <- sqrt(mean((pred- test2$wp2)^2))
mape <- mean(abs((test2$wp2 - pred)/test2$wp2))*100</pre>
```

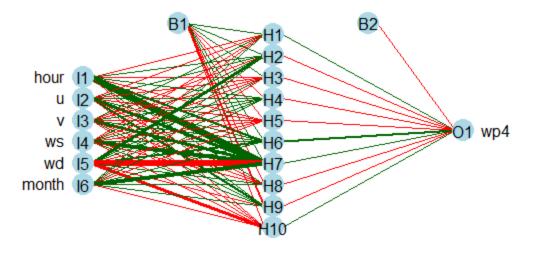




```
fml<- as.formula("wp4~.-Date-OnlyDate-year-month");
res <- nnet(fml, data=model4,size=10, linout=TRUE, skip=TRUE, MaxNWts=10000, trace=FALSE, maxit=1
summary(res)

#plot
plotnet(res,circle_col = "lightblue", circle_cex = 3,pos_col = "darkgreen", neg_col = "red")

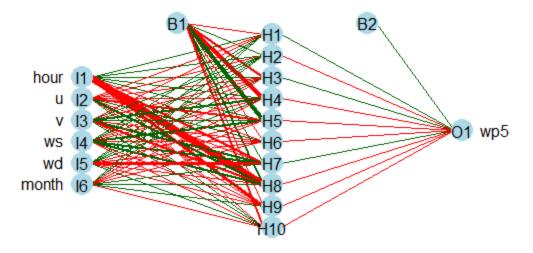
#predict
pred<-predict(res, newdata=test4,type="raw")
pred <- round(pred, digits = 3)
rmse <- sqrt(mean((pred- test4$wp4)^2))
mape <- mean(abs((test4$wp4 - pred)/test4$wp4))*100
rmse</pre>
```



```
fml<- as.formula("wp5~.-Date-OnlyDate-year");
res <- nnet(fml, data=model5,size=10, linout=TRUE, skip=TRUE, MaxNWts=10000, trace=FALSE, maxit=1
summary(res)

#plot
plotnet(res,circle_col = "lightblue", circle_kex = 3,pos_col = "darkgreen", neg_col = "red")

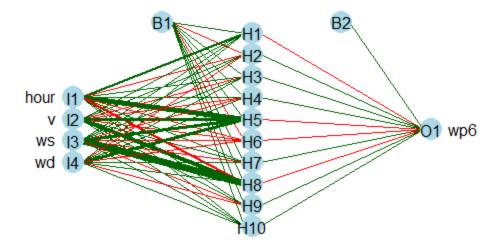
#predict
pred<-predict(res, newdata=test5,type="raw")
pred <- round(pred, digits = 3)
rmse <- sqrt(mean((pred- test5$wp5)^2))
mape <- mean(abs((test5$wp5 - pred)/test5$wp5))*100
rmse</pre>
```



```
fml<- as.formula("wp6~.-Date-OnlyDate-year-month-u");
res <- nnet(fml, data=model6,size=10, linout=TRUE, skip=TRUE, MaxNWts=10000, trace=FALSE, maxit=100)
summary(res)

#plot
plotnet(res,circle_col = "lightblue", circle_cex = 3,pos_col = "darkgreen", neg_col = "red")

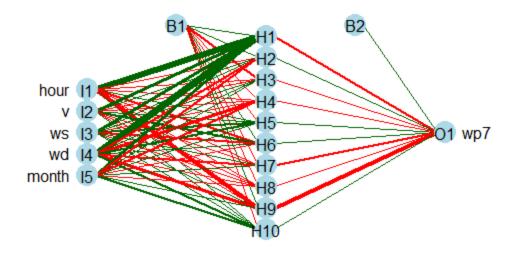
#predict
pred<-predict(res, newdata=test6,type="raw")
pred <- round(pred, digits = 3)
rmse <- sqrt[mean((pred- test6$wp6)^2))|
mape <- mean(abs((test6$wp6 - pred)/test6$wp6))*100
rmse</pre>
```



```
fml<- as.formula("wp7~.-Date-OnlyDate-year-u");
res <- nnet(fml, data=model7,size=10, linout=TRUE, skip=TRUE, MaxNWts=10000, trace=FALSE, maxit=100)
summary(res)

#plot
plotnet(res,circle_col = "lightblue", circle_cex = 3,pos_col = "darkgreen", neg_col = "red")

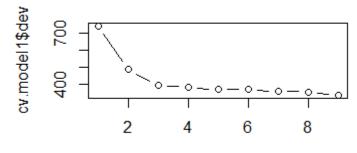
#predict
pred<-predict(res, newdata=test7,type="raw")
#pred <- round(pred, digits = 3)
rmse <- sqrt(mean((pred- test7$wp7)^2))
mape <- mean(abs((test7$wp7 - pred)/test7$wp7))*100
rmse</pre>
```



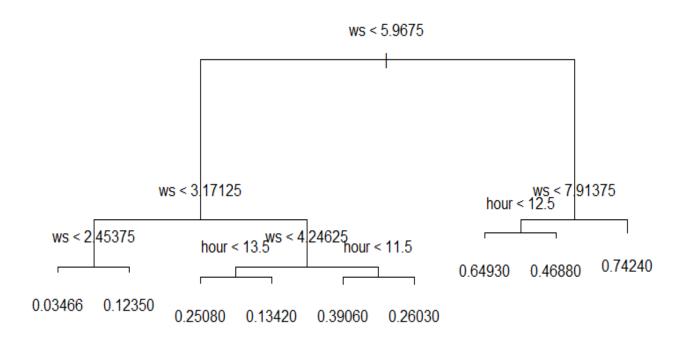
♣Regression Tree

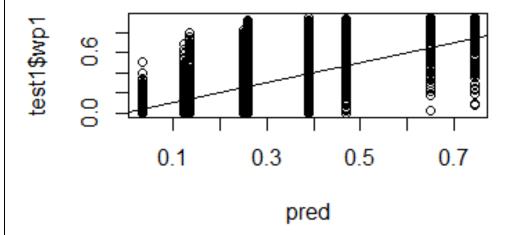
- 1. To go ahead with regression model we first cleaned the data and removed outliers with cook's method.
- 2. In regression tree we did modeling for all 7 farms and the models for all the farms were different.
- 3. Feature Selection was done all the farms and model was applied based on the different predictors.
- 4. Performance metrics and accuracy was calculated for all seven models.

```
library(tree)
tree.model1 = tree(wp1~.-Date-OnlyDate,model1)
summary (tree.model1)
#plot
plot (tree.model1)
text (tree.model1, pretty = 0)
cv.model1 = cv.tree (tree.model1)
plot (cv.model1$size, cv.model1$dev, type='b')
#predict
pred=predict (tree.model1, newdata =test1)
plot(pred,test1$wp1)
abline (0,1)
accuracy(pred,test1$wp1)
> summary (tree.model1)
Regression tree:
tree(formula = wp1 ~ . - Date - OnlyDate, data = model1)
Variables actually used in tree construction:
           "hour"
[1] "ws"
Number of terminal nodes: 9
Residual mean deviance: 0.02474 = 324.7 / 13120
Distribution of residuals:
         1st Qu.
    Min.
                     Median
                                 Mean 3rd Qu.
                                                    Max.
-0.63740 -0.10280 -0.02917 0.00000 0.08458 0.70780
> accuracy(pred,test1$wp1)
                  ME
                           RMSE
                                      MAE MPE MAPE
Test set 0.03795856 0.1835736 0.1358228 -Inf Inf
```

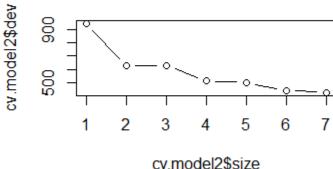


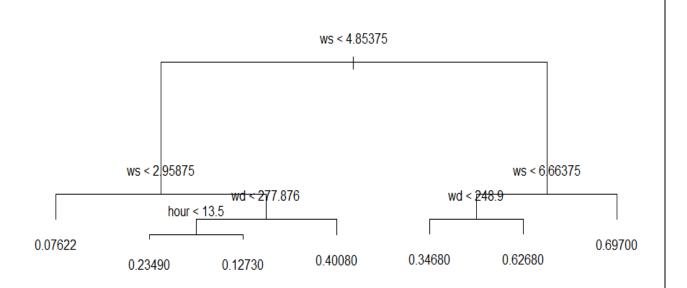
cv.model1\$size

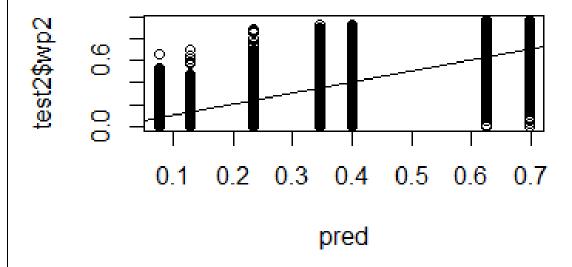




```
library(tree)
tree.model2 = tree(wp2~.-Date-OnlyDate,model2)
summary (tree.model2)
#plot
plot (tree.model2)
text (tree.model2, pretty = 0)
cv.model2 = cv.tree (tree.model2)
plot (cv.model2$size, cv.model2$dev, type='b')
#predict
pred=predict (tree.model2, newdata =test2)
plot(pred,test2$wp2)
abline (0,1)
accuracy(pred,test2$wp2)
> summary (tree.model2)
Regression tree:
tree(formula = wp2 ~ . - Date - OnlyDate, data = model2)
Variables actually used in tree construction:
[1] "ws" "wd"
                  "hour"
Number of terminal nodes: 7
Residual mean deviance: 0.03208 = 419 / 13060
Distribution of residuals:
    Min. 1st Qu.
                     Median
                                Mean 3rd Qu.
-0.63900 -0.08776 -0.03726 0.00000 0.09103 0.76480
> accuracy(pred,test2$wp2)
                            RMSE
                                       MAE MPE MAPE
                   ME
Test set -0.01006036 0.1774975 0.1347039 -Inf Inf
>
```







```
#RegressionTree
library(tree)
tree.model3 = tree(wp3~.-Date-OnlyDate,model3)
summary (tree.model3)

#plot
plot (tree.model3)
text (tree.model3, pretty = 0)
cv.model3 = cv.tree (tree.model3)
plot (cv.model3$size, cv.model3$dev, type='b')

#predict
pred=predict (tree.model3, newdata =test3)
plot(pred,test3$wp3)
abline (0,1)
accuracy(pred,test3$wp3)
```

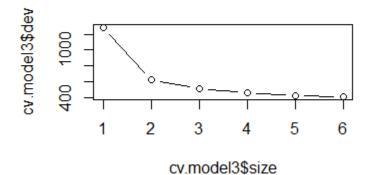
```
> summary (tree.model3)

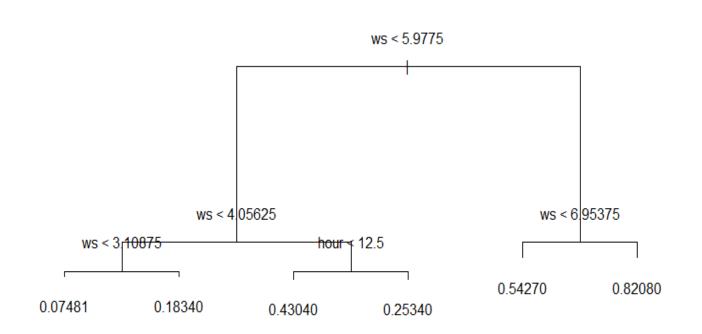
Regression tree:
tree(formula = wp3 ~ . - Date - OnlyDate, data = model3)
Variables actually used in tree construction:
[1] "ws" "hour"
Number of terminal nodes: 6
Residual mean deviance: 0.03109 = 407.2 / 13100
Distribution of residuals:
    Min. 1st Qu. Median Mean 3rd Qu. Max.
-0.82080 -0.08545 -0.02581 0.00000 0.10320 0.81420
```

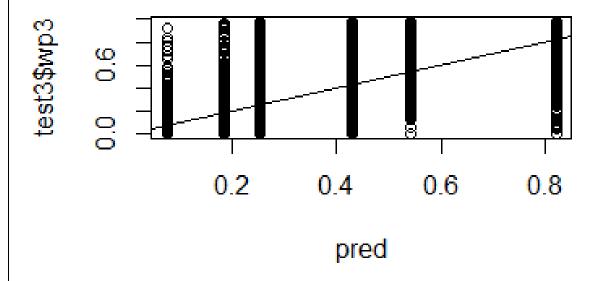
```
> accuracy(pred,test3$wp3)

ME RMSE MAE MPE MAPE

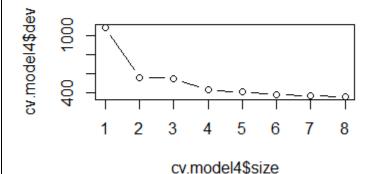
Test set 0.02290298 0.2005844 0.1452925 -Inf Inf
```

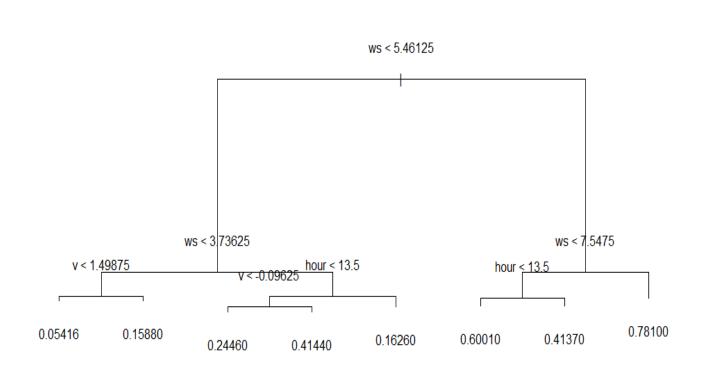


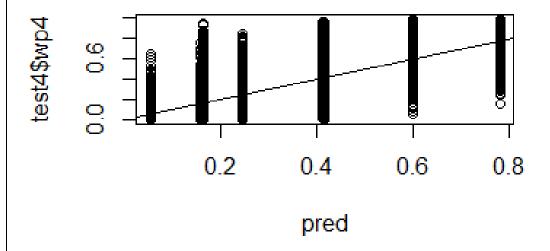




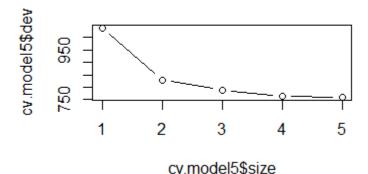
```
#RegressionTree
library(tree)
tree.model4 = tree(wp4~.-Date-OnlyDate,model4)
summary (tree.model4)
plot (tree.model4)
text (tree.model4, pretty = 0)
#predict
pred=predict (tree.model4, newdata =test4)
plot(pred,test4$wp4)
abline (0,1)
accuracy(pred,test4$wp4)
> summary (tree.mode14)
Regression tree:
tree(formula = wp4 ~ . - Date - OnlyDate, data = model4)
Variables actually used in tree construction:
                   "hour"
Number of terminal nodes: 8
Residual mean deviance: 0.02583 = 338.9 / 13120
Distribution of residuals:
    Min. 1st Qu.
                     Median
                                 Mean 3rd Qu.
-0.74200 -0.07824 -0.03369 0.00000 0.08884 0.70620
> accuracy(pred,test4$wp4)
                          RMSE
                                      MAE MPE MAPE
                  ME
Test set 0.02800967 0.1789393 0.1339365 -Inf Inf
```

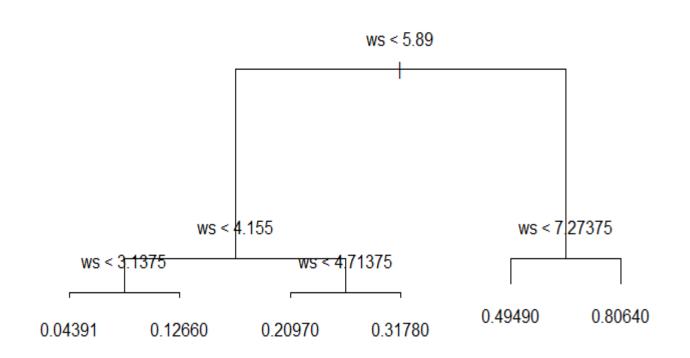


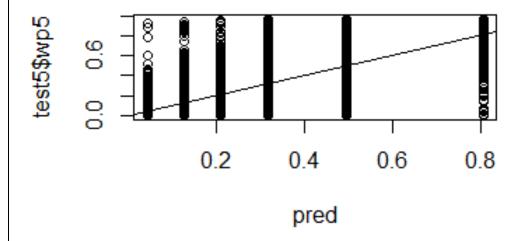




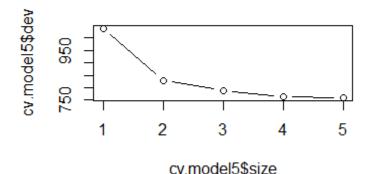
```
#RegressionTree
 library(tree)
 tree.model5 = tree(wp5~.-Date-OnlyDate,model5)
 summary (tree.model5)
 #plot
 plot (tree.model5)
 text (tree.model5, pretty = 0)
 #predict
 pred=predict (tree.model5, newdata =test5)
 plot(pred,test5$wp5)
 abline (0,1)
 accuracy(pred,test5$wp5)
> summary (tree.model5)
Regression tree:
tree(formula = wp5 ~ . - Date - OnlyDate, data = model5)
Variables actually used in tree construction:
[1] "ws"
Number of terminal nodes: 6
Residual mean deviance: 0.02631 = 343.9 / 13070
Distribution of residuals:
    Min. 1st Qu.
                     Median
                                 Mean 3rd Qu.
                                                    Max.
-0.80640 -0.08375 -0.02891 0.00000 0.07618 0.77840
> accuracy(pred,test5$wp5)
                 ME
                         RMSE
                                     MAE MPE MAPE
Test set 0.0485396 0.1979785 0.1387474 -Inf Inf
```

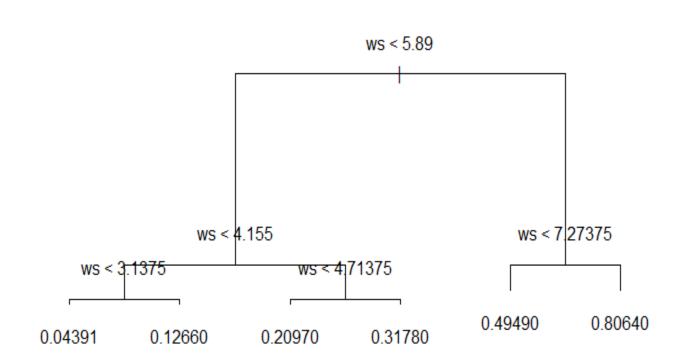


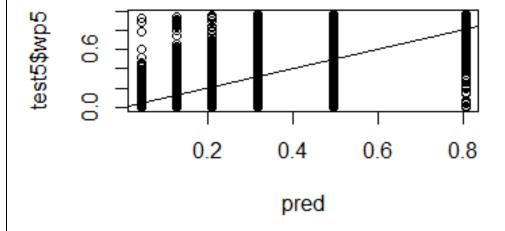




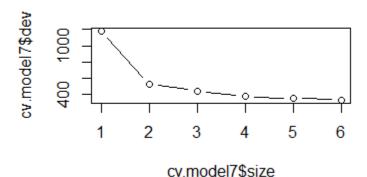
```
#RegressionTree
 library(tree)
 tree.model6 = tree(wp6~.-Date-OnlyDate,model6)
summary (tree.model6)
#plot
plot (tree.model6)
text (tree.model6, pretty = 0)
#predict
pred=predict (tree.model6, newdata =test6)
plot(pred,test6$wp6)
abline (0,1)
accuracy(pred,test6$wp6)
> summary (tree.model5)
Regression tree:
tree(formula = wp5 ~ . - Date - OnlyDate, data = model5)
Variables actually used in tree construction:
[1] "ws"
Number of terminal nodes: 6
Residual mean deviance: 0.02631 = 343.9 / 13070
Distribution of residuals:
    Min. 1st Qu.
                    Median
                                 Mean 3rd Qu.
                                                    Max.
-0.80640 -0.08375 -0.02891 0.00000 0.07618 0.77840
> accuracy(pred,test5$wp5)
                 ME
                         RMSE
                                     MAE MPE MAPE
Test set 0.0485396 0.1979785 0.1387474 -Inf Inf
>
```

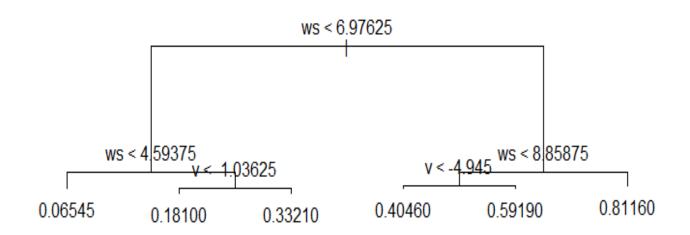


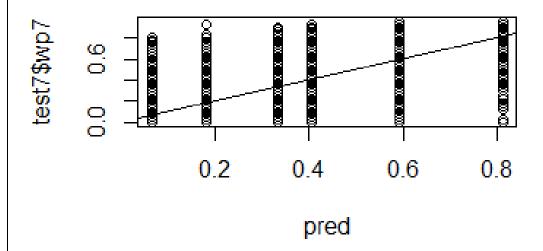




```
#RegressionTree
library(tree)
tree.model7 = tree(wp7~.-Date-OnlyDate,model7)
summary (tree.model7)
#plot
plot (tree.model7)
text (tree.model7, pretty = 0)
cv.model7 = cv.tree (tree.model7)
plot (cv.model7$size, cv.model7$dev, type='b')
#predict
pred=predict (tree.model7, newdata =test7)
plot(pred,test7$wp7)
abline (0,1)
accuracy(pred,test7$wp7)
> summary (tree.model7)
Regression tree:
tree(formula = wp7 ~ . - Date - OnlyDate, data = model7)
Variables actually used in tree construction:
[1] "ws" "v"
Number of terminal nodes: 6
Residual mean deviance: 0.02495 = 326.6 / 13090
Distribution of residuals:
    Min. 1st Qu.
                     Median
                                 Mean 3rd Qu.
                                                    мах.
-0.71060 -0.06545 -0.04045 0.00000 0.09687 0.70300
> plot (tree.model7)
> accuracy(pred,test7$wp7)
                          RMSE
                                     MAE MPE MAPE
                 ME
Test set 0.0117528 0.1671721 0.1240614 -Inf Inf
>
```







Discussion

- ✓ With the results derived above we can conclude that Multi Linear Regression has good Performance metrics.
- ✓ The forecast values had good varied predictions as compared to decision tree and neural network as these both divide the region and give a fixed values in when a data falls in a particular region
- ✓ Multilinear Regression works well in case of regression while neural network and decision tree works best in case of classification.

Benchmark Output files: -







benchmark_nnet.csv

benchmark_tree.csv

benchmark_linearmodel.csv
